

Presentation Transcript & Speaker Notes

Twenty slides • 30:00 total run time • full narration and source cues.

THE INTELLIGENCE LAYER FOR MARS

Release 5 • 04 September 2026

Presenter transcript and notes

Designed for a 30-minute convention briefing. Timings are embedded in the source deck; speaker notes preserve claims, transitions, and source cues.

01 • 29TH INTERNATIONAL MARS SOCIETY CONVENTION

[Timing: 1:00] Mars is not waiting for one discipline at a time. Every meaningful surface system couples science, physical engineering, operations and human survival. The thesis tonight is simple: the next leap is not a larger chatbot. It is an intelligence layer—a governed way for people, AI models, solvers, instruments and robots to work through evidence, design, verification and learning as one system. The harness is the durable part. Models will change rapidly. Mars missions cannot rebuild their operating method every time the model underneath improves. We need persistent memory, explicit work products, traceable handoffs and gates that stronger future intelligence can inherit. The goal is faster learning without counterfeit certainty. [Sources] -

<https://www.marssociety.org/news/2026/03/23/save-the-dates-29th-annual-international-mars-society-convention-at-usc/> - <https://convention.marssociety.org/> - Bob J Mills, Detailed Initial Chat Summary and Restored Concepts, Mars project record. - Visual: original AI-generated Mars settlement image created for this deck.

02 • 02 / MISSION CONSTRAINT

[Timing: 1:15] Frame the problem as learning latency, not merely compute. Mars teams must turn uncertain observations into defensible action while mass, power, schedule and crew margins keep closing. That requires literature search, data grounding, hypothesis testing, multidisciplinary analysis, review, documentation and physical verification to move as a connected loop. Communications delay makes the problem visible. NASA's Moon-to-Mars analysis finds one-way delays can reach roughly twenty-one to twenty-three minutes, with representative blackout periods lasting about thirteen to twenty-one days. Earth cannot remain the synchronous brain for every decision. The same constraint appears before launch: science and engineering teams lose time at handoffs, review preparation and rediscovery of prior reasoning. Agentic harnesses do not make physics instantaneous. They compress avoidable waiting by running controlled work in parallel, carrying evidence forward and surfacing broken links early. [Sources] -

<https://www.nasa.gov/wp-content/uploads/2024/01/mars-communications-disruption-and-delay.pdf> -

<https://science.nasa.gov/mars/facts/> - Bob J Mills, Version 4 transcript and restored presentation design record. - Visual: original AI-generated Mars mission-control image created for this deck.

03 • 03 / DESIGN RULE

[Timing: 1:15] This is the architectural definition. A model is the engine; the harness is the infrastructure surrounding it. That infrastructure is the difference between describing a task and completing it with evidence and accountable control. Walk the seven functions: first, the orchestration loop that plans, executes, inspects and repairs. Second, system instructions that define role and purpose. Third, tool access. Fourth, context management. Fifth, memory and state. Sixth, bounded subagents. Seventh, guardrails and approval controls. The model can change without breaking the process because work-product contracts, ground-truth checks, durable state and authority remain stable. This is why harness design can matter more than the model itself. [Sources] - User-supplied agent-harness definition and seven-function brief.

04 • 04 / CAPABILITY MAP

[Timing: 1:15] Named agents are useful for implementation, but they are a fragile architecture. The durable map is the set of capabilities and the evidence each capability must produce. Scientific reasoning, literature grounding, simulation, systems engineering, CAD, FMEA, test planning and governance are skills. A person, model, database, solver, robot or future AGI can perform a skill if it satisfies the contract. This matters on Mars because the executor changes with context. During communications delay, a local model may act. For a nonlinear trade, a deterministic solver should act. For a physical claim, an instrument or chamber test must act. For risk acceptance, authorized mission leadership must act. Skills over agents prevents tool worship. It also makes the handoff auditable: the next stage receives an approved work product, not merely the opinion of a named digital persona. [Sources] - Bob J Mills, Detailed Initial Chat Summary and Restored Concepts, section 5. - <https://github.com/Robertstar2000/Hypatia-Pro> - <https://github.com/Robertstar2000/Intelligent-Engineer>

05 • 05 / SYSTEM ARCHITECTURE

[Timing: 1:15] The architecture uses two connected harnesses because science and engineering have different jobs. Hypatia Pro asks what is supported, uncertain or unresolved. Intelligent Engineer Pro asks how to build within that evidence, how to control risk and how to prove the result. The forward handoff is an Engineering Evidence Package: equations, distributions, operating envelopes, provenance, assumptions, failure thresholds and minimum additional tests. The return handoff is a Research Request Package: a blocked decision, competing hypotheses, variables, Mars boundary conditions, method and acceptance threshold. That return path is essential. It turns “we need more research” into a bounded job that can be executed and reinserted into the same baseline. Verification, manufacturing and operations then generate new evidence. The digital thread keeps the loop continuous across tools, people and future model generations. [Sources] - Bob J Mills, Detailed Initial Chat Summary and Restored Concepts, sections 6–8. - <https://github.com/Robertstar2000/Hypatia-Pro> - <https://github.com/Robertstar2000/Intelligent-Engineer> - Animated telemetry assets: original AI-generated sprite sheet assembled into GIFs for this deck.

06 • 06 / SCIENCE HARNESS

[Timing: 1:20] Hypatia is a scientific-discovery harness, not a fluent literature summarizer. It refines the question, grounds it in prior evidence, generates falsifiable competing hypotheses, defines variables and measurement protocols, obtains evidence, quantifies uncertainty and invites adversarial attack. The governing loop is generate, validate, critique, repair, approve and persist. A model may draft an artifact, but drafting is not promotion. The artifact must survive checks appropriate to its class. A literature claim needs a traceable source. A simulated result needs a model, assumptions and boundary conditions. An empirical claim needs measurement provenance. A failed hypothesis is preserved because negative evidence is part of mission knowledge. The loop can move backward whenever grounding, protocol or inference is weak. That willingness to refuse premature certainty is what makes the output safe enough to become an engineering input. [Sources] - Bob J Mills, Detailed Initial Chat Summary and Restored Concepts, section 6. - <https://github.com/Robertstar2000/Hypatia-Pro> - Visual: original AI-generated Mars laboratory image created for this deck.

07 • 07 / EVIDENCE CONTROL

[Timing: 1:20] A mission loses integrity when evidence classes blur. Literature can be authoritative within its scope, but it is not automatically valid at a different pressure, temperature, material or scale. A model can be useful and still be wrong outside its assumptions. A physical measurement can be real and still have uncertainty or limited transferability. An unresolved question is not a defect to hide; it is a work item. The harness attaches status to the claim itself. It also carries source, scope, assumptions, boundary conditions, measurement method, uncertainty and disposition. That allows engineering to use knowledge without pretending it is stronger than it is. This is especially important in AI-generated work. Fluency can erase the visual cues that used to warn us an answer was tentative. Explicit epistemic status puts those cues back into the engineering record and lets later evidence reopen affected requirements and decisions. [Sources] - Bob J Mills, Sources, Claims and Replacement Ledger, Mars Version 4 package. - <https://science.nasa.gov/mars/facts/> - <https://github.com/Robertstar2000/Hypatia-Pro>

08 • 08 / ENGINEERING HARNESS

[Timing: 1:20] Intelligent Engineer Pro takes an approved evidence package and converts it into requirements, architecture alternatives, margins, design analyses, risks, verification methods and controlled review artifacts. It is phase-gated because a fast draft must not become an invisible baseline. The harness supports project definition, requirements and constraints, discipline selection, preliminary and critical design, focused design sprints, multidisciplinary analysis, CAD, FEA, CFD, circuit and process work, DFMA, FMEA, redundancy, manufacturing, maintenance and operations planning. It can generate next-phase seeds, but those seeds remain drafts until reviewed. The important distinction is authorization. A model can propose a requirement or design. The project records what evidence supports it, what assumptions remain, what test will verify it and what authority locks it. That discipline lets automation accelerate engineering without smuggling unreviewed content into the next phase. [Sources] - Bob J Mills, Detailed Initial Chat Summary and Restored Concepts, section 6. - <https://github.com/Robertstar2000/Intelligent-Engineer> - Visual: original AI-generated Mars engineering and digital-twin image created for this deck.

09 • 09 / DIGITAL THREAD

[Timing: 1:20] The handoff between science and engineering cannot be a prose answer. Engineering needs a structured evidence package containing equations, parameter distributions, operating envelopes, assumptions, failure thresholds, confidence, provenance and the tests still required. Each element is linked. A source supports a bounded claim. The claim defines a parameter or constraint. The parameter supports a requirement or trade. The requirement has a verification method and status. When evidence changes, the digital thread identifies exactly which requirement, model, test or decision is affected. This is where a harness becomes more than workflow automation. It creates machine-readable contracts between disciplines. It allows agents to maintain traceability, detect inconsistencies and reopen focused work before a review. The result is less time reconstructing intent and more time resolving the actual engineering question. [Sources] - Bob J Mills, Detailed Initial Chat Summary and Restored Concepts, section 7. - <https://github.com/Robertstar2000/Hypatia-Pro> - <https://github.com/Robertstar2000/Intelligent-Engineer>

10 • 10 / COMPLETION LOOP

[Timing: 1:20] A linear research-to-engineering arrow is incomplete. The real acceleration comes from closing the return path. Research delivers approved evidence. Engineering turns it into requirements and trades. Design or analysis exposes a blocking uncertainty. The harness packages that uncertainty as a precise research request. Targeted evidence returns to the same baseline, the affected decision is updated and engineering continues. Agents can search, simulate, critique, document and maintain links throughout the loop. The key is that they do not erase the issue. They make it explicit and executable. A good Research Request Package names the blocked decision, hypotheses, variables, Mars boundary conditions, method, minimum detectable effect, evidence quality and acceptance threshold. Verification, manufacturing and operations remain part of the same learning system. A failed test or field observation can reopen science, requirements or design without losing the chain of reasoning. [Sources] - Bob J Mills, Detailed Initial Chat Summary and Restored Concepts, section 8. - <https://github.com/Robertstar2000/Hypatia-Pro> - <https://github.com/Robertstar2000/Intelligent-Engineer>

11 • 11 / AGENTIC ACCELERATION

[Timing: 1:10] Parallelism is where agentic systems can attack real schedule. One workstream searches evidence and contradictions. Another runs models and sensitivity studies. Another tries to break the preferred conclusion. Another builds traceable artifacts and review material. These tasks can overlap when their interfaces and permissions are defined. The merge gate matters more than the number of agents. Independent outputs must be reconciled, validated, critiqued and approved before they become one package. Parallel work should not create four incompatible truths or hide duplicated assumptions. The harness records dependencies, evidence classes and unresolved conflicts. The result is not “AI worked four times faster.” The defensible claim is narrower: controlled concurrency reduces avoidable waiting and review preparation while preserving the evidence standard. That is a practical way to compress the research–engineering–research issue–completion loop. [Sources] - Bob J Mills, Version 4 Change Log and restored presentation design record. - <https://github.com/Robertstar2000/Hypatia-Pro> - <https://github.com/Robertstar2000/Intelligent-Engineer>

12 • 12 / VERIFICATION

[Timing: 1:10] Verification is not the last phase. It begins with the first claim. Every claim needs provenance. Every requirement needs a method. Every uncertainty needs a disposition: analyze it, test it, constrain the design around it or explicitly accept the risk. A harness can maintain the trace matrix continuously. It can identify a source that no longer supports its claim, a changed parameter that invalidates a requirement, a requirement without a method or a test result that does not meet its acceptance criterion. That broken link becomes focused rework before the review, not a surprise during it. AI can help generate tests, check consistency, update FMEA and organize evidence. It cannot claim that a simulation physically validated hardware. Verification is the discipline that separates an impressive demonstration from engineering that can be trusted on Mars. [Sources] - Bob J Mills, Detailed Initial Chat Summary and Restored Concepts, section 9. - <https://ntrs.nasa.gov/citations/20100036405> - <https://github.com/Robertstar2000/Intelligent-Engineer>

13 • 13 / MISSION MEMORY

[Timing: 1:10] The shared digital brain is not a personified super-AI. It is governed mission context: evidence, assumptions, parameters, requirements, decisions, risks, tests, field data and open questions connected in a reusable graph. That memory survives staff changes, model upgrades and tool replacements. It prevents context from disappearing into isolated chats or static reports. It also makes future models more useful because they inherit structured knowledge instead of reconstructing

the mission from fragments. Every harness reads and writes through controlled interfaces. Hypatia adds approved evidence and uncertainty. Intelligent Engineer adds requirements, trades and verification. CAD, PLM, simulators, labs and operations add models and results. The digital thread links them. This is how learning compounds across projects: a negative test, failed hypothesis or field observation remains available to prevent repetition and sharpen the next decision. [Sources] - Bob J Mills, Detailed Initial Chat Summary and Restored Concepts, section 11. - <https://github.com/Robertstar2000/Hypatia-Pro> - <https://github.com/Robertstar2000/Intelligent-Engineer>

14 • 14 / MARS AUTONOMY

[Timing: 1:10] NASA's communication analysis makes the autonomy requirement unavoidable. One-way delay can exceed twenty minutes, and solar conjunction can produce blackouts lasting roughly two or three weeks in representative mission profiles. Relays may mitigate some disruption, but they cannot eliminate the speed-of-light delay. A Mars intelligence layer therefore needs three operating modes. Local mode uses cached evidence, local models and controlled tools to keep work moving. Sync mode reconciles new evidence, decisions and software with Earth when the link is available. Safe mode degrades predictably when capabilities, data or authority are insufficient. The harness makes those boundaries explicit. It knows which decisions can be made locally, which require later reconciliation and which must stop. This is more robust than treating autonomy as a single "agent" and more realistic than assuming mission control remains continuously synchronous. [Sources] - <https://www.nasa.gov/wp-content/uploads/2024/01/mars-communications-disruption-and-delay.pdf> - <https://science.nasa.gov/mars/facts/> - Visual: original AI-generated Earth–Mars communications image created for this deck.

15 • 15 / CASE STUDY — SCIENCE OUTPUT

[Timing: 2:30] Now apply the architecture to a compact Mars oxygen plant. NASA's MOXIE experiment gives us a real empirical anchor. Across sixteen runs it produced a total of 122 grams of oxygen. At peak it reached 12 grams per hour—twice its original goal—at 98 percent purity or better. That is powerful evidence that solid-oxide electrolysis can extract oxygen from the Martian carbon-dioxide atmosphere. But a settlement plant is not simply "make MOXIE bigger." NASA's own next-step framing includes the generator plus oxygen liquefaction and storage. Engineering must also resolve atmospheric variability, dust filtration and electrostatics, thermal cycling, material and seal degradation, process aging, contamination, maintenance access, redundancy and interfaces to power and storage. A Hypatia evidence package would separate the empirical MOXIE results from literature, modeled expectations and open questions. It would carry equations, parameter ranges, boundary conditions, uncertainty, failed hypotheses and the minimum additional experiments required. The visible open list is intentionally not presented as solved. This is the integrity point: proof of principle is evidence, but it is not yet a settlement-qualified baseline. [Sources] - <https://www.nasa.gov/missions/mars-2020-perseverance/perseverance-rover/nasa-s-oxygen-generating-experiment-moxie-completes-mars-mission/> - <https://science.nasa.gov/mars/facts/> - <https://ntrs.nasa.gov/citations/20100036405> - Bob J Mills, Hypatia / Intelligent Engineer Demo Package, Mars Version 4. - Visual: original AI-generated Mars ISRU cutaway created for this deck.

16 • 16 / CASE STUDY — ENGINEERING OUTPUT

[Timing: 2:30] This slide demonstrates how Intelligent Engineer should translate evidence into a controlled baseline. It is explicitly illustrative; the final conference version should replace these rows with reviewed harness output. The first track begins with MOXIE's empirical proof of principle. It does not invent a settlement production number. It creates a requirement to meet mission-defined oxygen demand with qualified margin, verified by an integrated endurance demonstration. The second track uses the thin carbon-dioxide atmosphere and the need for an inlet envelope. The requirement must reference the qualified pressure and composition profile, then be verified in a representative chamber. The third track refuses to pretend the dust issue is solved. It creates a recoverability requirement tied to a specified exposure, with a dust loading and recovery test. The fourth track follows NASA's stated future direction: a full system needs oxygen generation plus liquefaction and storage, so engineering defines and verifies that interface. Notice the pattern. Evidence may support a requirement. Open uncertainty may create a bounded assumption, margin or research request. A mission decision may deliberately set a value. Every path remains visible. Every requirement has a proof method from the moment it enters the baseline. [Sources] - <https://www.nasa.gov/missions/mars-2020-perseverance/perseverance-rover/nasas-oxygen-generating-experiment-moxie-completes-mars-mission/> - <https://science.nasa.gov/mars/facts/> - <https://ntrs.nasa.gov/citations/20100036405> - Bob J Mills, Hypatia / Intelligent Engineer Demo Package, Mars Version 4. - <https://github.com/Robertstar2000/Intelligent-Engineer>

17 • 17 / CASE STUDY — RETURN PATH

[Timing: 2:30] Here the engineering analysis has reached a real decision: how often must the inlet filter be cleaned or replaced? That decision affects filter area, compressor margin, maintenance time, spares and operating availability. But dust adhesion after repeated thermal cycling under Mars pressure is not sufficiently bounded. Instead of burying a convenient assumption, Intelligent Engineer creates a Research Request Package. It states competing hypotheses. One says thermal cycling has limited effect within the qualified range. Another says electrostatic adhesion drives persistent fouling. It defines the dependent measures—pressure drop, recovery and adhesion—and the independent variables—dust loading, temperature cycles and exposure. It specifies Mars-relevant pressure and composition, a qualified simulant, the proposed factorial chamber test and uncertainty analysis. The acceptance criterion is not invented by the scientist; it is derived from the engineering margin and requirement. Hypatia executes or coordinates the targeted work, classifies the evidence, attacks the inference and returns only reviewed findings. The baseline then updates the filter sizing, compressor margin, cleaning interval and verification plan. Engineering continues from the same traceable state. This is the completion loop the presentation is advocating. [Sources] - <https://ntrs.nasa.gov/citations/20100036405> - <https://www.nasa.gov/science-research/kennedy-scientists-developing-technology-to-remove-martian-dust/> - <https://science.nasa.gov/mars/facts/> - Bob J Mills, Hypatia / Intelligent Engineer Demo Package and restored completion-loop requirement. - <https://github.com/Robertstar2000/Hypatia-Pro> - <https://github.com/Robertstar2000/Intelligent-Engineer>

18 • 18 / HARNESS ECOSYSTEM

[Timing: 1:50] The five application patterns become portable skills. Hypatia structures scientific work. Hyperia performs broad research and synthesis. Intelligent Engineer controls the physical-engineering baseline. Vibe Engineering supports collaborative design. PM Accelerator controls execution and review readiness. Each pack defines a role, decision loop, tool contract, work product, ground-truth checks, memory rules and approval boundaries. Three shared controls—orchestration, validation and mission memory—compose them safely. The core instructions can travel across Codex, Claude Cowork, Hermes and Grok. Placement and invocation differ, and the runtime can execute only the tools and permissions it actually has. This is functional workflow emulation, not a copy of a proprietary interface or a claim of feature parity. [Sources] - <https://docs.x.ai/build/features/skills-plugins-marketplaces> - <https://github.com/NousResearch/hermes-agent/blob/main/website/docs/user-guide/features/skills.md> - <https://support.claude.com/en/articles/12512198-how-to-create-custom-skills> - User-supplied SaaS application concepts.

19 • 19 / ADOPTION

[Timing: 1:40] The practical starting point is not an enterprise transformation. Choose one recurring decision where delay, rework or lost evidence matters. Define the required work products, evidence classes, tools, permissions, validation steps and release authority. Connect only the data and tools needed for that loop. Then compare the harness output against expert work and known failure cases. Can it preserve provenance? Does it distinguish modeled from empirical evidence? Does it expose uncertainty? Can a reviewer reconstruct the decision? Does it fail safely when the evidence is inadequate? Only after that loop survives real pressure should the organization scale to adjacent skills and domains. Most organizations are still moving from chatbot to task agent. The practical next stage is the harness: workflow, memory, tools and gates. Mars-scale learning eventually requires an ecosystem, but reliability is earned one consequential loop at a time. [Sources] - Bob J Mills, Version 4 Run of Show and restored adoption model. - <https://github.com/Robertstar2000/Hypatia-Pro> - <https://github.com/Robertstar2000/Intelligent-Engineer>

20 • 20 / CALL TO ACTION

[Timing: 1:40] Return to the opening promise. Mars is a learning-rate problem, but speed without discipline is not progress. The architecture presented here turns AI into structured engineering capacity: specialized harnesses, durable skills, explicit evidence classes, phase gates, continuous verification, persistent mission memory and a complete return path from engineering uncertainty to targeted science. The model is replaceable. The way the mission defines work, remembers evidence, accepts risk and proves a result must endure. Stronger AI will move more execution into automation. That makes the harness more important, not less, because purpose, authority, provenance and acceptance must remain visible as capability changes. Start with one consequential loop. Make the evidence survive review. Preserve negative results. Let science refuse false certainty. Let engineering return precise questions. Then connect the loops until science, engineering,

verification and operations become one learning system. Mars will not be built by a chatbot. It will be built by people and increasingly capable intelligence working through a mission-ready way to learn. [Sources] - Bob J Mills, Detailed Initial Chat Summary and Restored Concepts, durable principles and closing frame. - <https://www.nasa.gov/wp-content/uploads/2024/01/mars-communications-disruption-and-delay.pdf> - <https://www.nasa.gov/missions/mars-2020-perseverance/perseverance-rover/nasas-oxygen-generating-experiment-moxie-completes-mars-mission/> - <https://github.com/Robertstar2000/Hypatia-Pro> - <https://github.com/Robertstar2000/Intelligent-Engineer> - Visual: original AI-generated closing Mars settlement image created for this deck.

Apache License 2.0

Copyright 2026 Mars Harness Lab contributors

<https://www.apache.org/licenses/LICENSE-2.0>

Apache License
Version 2.0, January 2004
<http://www.apache.org/licenses/>

TERMS AND CONDITIONS FOR USE, REPRODUCTION, AND DISTRIBUTION

1. Definitions.

"License" shall mean the terms and conditions for use, reproduction, and distribution as defined by Sections 1 through 9 of this document.

"Licensor" shall mean the copyright owner or entity authorized by the copyright owner that is granting the License.

"Legal Entity" shall mean the union of the acting entity and all other entities that control, are controlled by, or are under common control with that entity. For the purposes of this definition, "control" means (i) the power, direct or indirect, to cause the direction or management of such entity, whether by contract or otherwise, or (ii) ownership of fifty percent (50%) or more of the outstanding shares, or (iii) beneficial ownership of such entity.

"You" (or "Your") shall mean an individual or Legal Entity exercising permissions granted by this License.

"Source" form shall mean the preferred form for making modifications, including but not limited to software source code, documentation source, and configuration files.

"Object" form shall mean any form resulting from mechanical transformation or translation of a Source form, including but not limited to compiled object code, generated documentation, and conversions to other media types.

"Work" shall mean the work of authorship, whether in Source or Object form, made available under the License, as indicated by a copyright notice that is included in or attached to the work (an example is provided in the Appendix below).

"Derivative Works" shall mean any work, whether in Source or Object form, that is based on (or derived from) the Work and for which the editorial revisions, annotations, elaborations, or other modifications represent, as a whole, an original work of authorship. For the purposes of this License, Derivative Works shall not include works that remain separable from, or merely link (or bind by name) to the interfaces of, the Work and Derivative Works thereof.

"Contribution" shall mean any work of authorship, including the original version of the Work and any modifications or additions to that Work or Derivative Works thereof, that is intentionally submitted to Licensor for inclusion in the Work by the copyright owner or by an individual or Legal Entity authorized to submit on behalf of the copyright owner. For the purposes of this definition, "submitted" means any form of electronic, verbal, or written communication sent to the Licensor or its representatives, including but not limited to communication on electronic mailing lists, source code control systems, and issue tracking systems that are managed by, or on behalf of, the Licensor for the purpose of discussing and improving the Work, but excluding communication that is conspicuously marked or otherwise designated in writing by the copyright owner as "Not a Contribution."

"Contributor" shall mean Licensor and any individual or Legal Entity on behalf of whom a Contribution has been received by Licensor and subsequently incorporated within the Work.

2. Grant of Copyright License. Subject to the terms and conditions of this License, each Contributor hereby grants to You a perpetual, worldwide, non-exclusive, no-charge, royalty-free, irrevocable copyright license to reproduce, prepare Derivative Works of, publicly display, publicly perform, sublicense, and distribute the Work and such Derivative Works in Source or Object form.
3. Grant of Patent License. Subject to the terms and conditions of this License, each Contributor hereby grants to You a perpetual, worldwide, non-exclusive, no-charge, royalty-free, irrevocable (except as stated in this section) patent license to make, have made, use, offer to sell, sell, import, and otherwise transfer the Work, where such license applies only to those patent claims licensable by such Contributor that are necessarily infringed by their Contribution(s) alone or by combination of their Contribution(s) with the Work to which such Contribution(s) was submitted. If You institute patent litigation against any entity (including a cross-claim or counterclaim in a lawsuit) alleging that the Work or a Contribution incorporated within the Work constitutes direct or contributory patent infringement, then any patent

Apache License 2.0

Copyright 2026 Mars Harness Lab contributors

<https://www.apache.org/licenses/LICENSE-2.0>

licenses
granted to You under this License for that Work shall
terminate
as of the date such litigation is filed.

4. Redistribution. You may reproduce and distribute copies of the

Work or Derivative Works thereof in any medium, with or without
modifications, and in Source or Object form, provided that You
meet the following conditions:

- (a) You must give any other recipients of the Work or Derivative Works a copy of this License; and
- (b) You must cause any modified files to carry prominent notices stating that You changed the files; and
- (c) You must retain, in the Source form of any Derivative Works that You distribute, all copyright, patent, trademark, and attribution notices from the Source form of the Work, excluding those notices that do not pertain to any part of the Derivative Works; and
- (d) If the Work includes a "NOTICE" text file as part of its distribution, then any Derivative Works that You distribute must include a readable copy of the attribution notices contained within such NOTICE file, excluding those notices that do not pertain to any part of the Derivative Works, in at least one of the following places: within a NOTICE text file distributed as part of the Derivative Works; within the Source form or documentation, if provided along with the Derivative Works; or, within a display generated by the Derivative Works, if and wherever such third-party notices normally appear. The contents of the NOTICE file are for informational purposes only and do not modify the License. You may add Your own attribution notices within Derivative Works that You distribute, alongside or as an addendum to the NOTICE text from the Work, provided that such additional attribution notices cannot be construed as modifying the License.

You may add Your own copyright statement to Your modifications and may provide additional or different license terms and conditions for use, reproduction, or distribution of Your modifications, or for any such Derivative Works as a whole, provided Your use, reproduction, and distribution of the Work otherwise complies with

the conditions stated in this License.

- 5. Submission of Contributions. Unless You explicitly state otherwise, any Contribution intentionally submitted for inclusion in the Work by You to the Licensor shall be under the terms and conditions of this License, without any additional terms or conditions. Notwithstanding the above, nothing herein shall supersede or modify the terms of any separate license agreement you may have executed with Licensor regarding such Contributions.

- 6. Trademarks. This License does not grant permission to use the trade names, trademarks, service marks, or product names of the Licensor, except as required for reasonable and customary use in describing the origin of the Work and reproducing the content of the NOTICE file.

- 7. Disclaimer of Warranty. Unless required by applicable law or agreed to in writing, Licensor provides the Work (and each Contributor provides its Contributions) on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied, including, without limitation, any warranties or conditions of TITLE, NON-INFRINGEMENT, MERCHANTABILITY, or FITNESS FOR A PARTICULAR PURPOSE. You are solely responsible for determining the appropriateness of using or redistributing the Work and assume any risks associated with Your exercise of permissions under this License.

- 8. Limitation of Liability. In no event and under no legal theory, whether in tort (including negligence), contract, or otherwise, unless required by applicable law (such as deliberate and grossly negligent acts) or agreed to in writing, shall any Contributor be liable to You for damages, including any direct, indirect, special, incidental, or consequential damages of any character arising as a result of this License or out of the use or inability to use the Work (including but not limited to damages for loss of goodwill, work stoppage, computer failure or malfunction, or any and all other commercial damages or losses), even if such Contributor has been advised of the possibility of such damages.

- 9. Accepting Warranty or Additional Liability. While redistributing the Work or Derivative Works thereof, You may choose to offer, and charge a fee for, acceptance of support, warranty, indemnity,

Apache License 2.0

Copyright 2026 Mars Harness Lab contributors

<https://www.apache.org/licenses/LICENSE-2.0>

or other liability obligations and/or rights consistent with this License. However, in accepting such obligations, You may act only on Your own behalf and on Your sole responsibility, not on behalf of any other Contributor, and only if You agree to indemnify, defend, and hold each Contributor harmless for any liability incurred by, or claims asserted against, such Contributor by reason of your accepting any such warranty or additional liability.

END OF TERMS AND CONDITIONS

APPENDIX: How to apply the Apache License to your work.

To apply the Apache License to your work, attach the following boilerplate notice, with the fields enclosed by brackets "[]" replaced with your own identifying information. (Don't include the brackets!) The text should be enclosed in the appropriate comment syntax for the file format. We also recommend that a file or class name and description of purpose be included on the same "printed page" as the copyright notice for easier identification within third-party archives.

Copyright [yyyy] [name of copyright owner]

Licensed under the Apache License, Version 2.0 (the "License");
you may not use this file except in compliance with the License.
You may obtain a copy of the License at

<http://www.apache.org/licenses/LICENSE-2.0>

Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
See the License for the specific language governing permissions and limitations under the License.